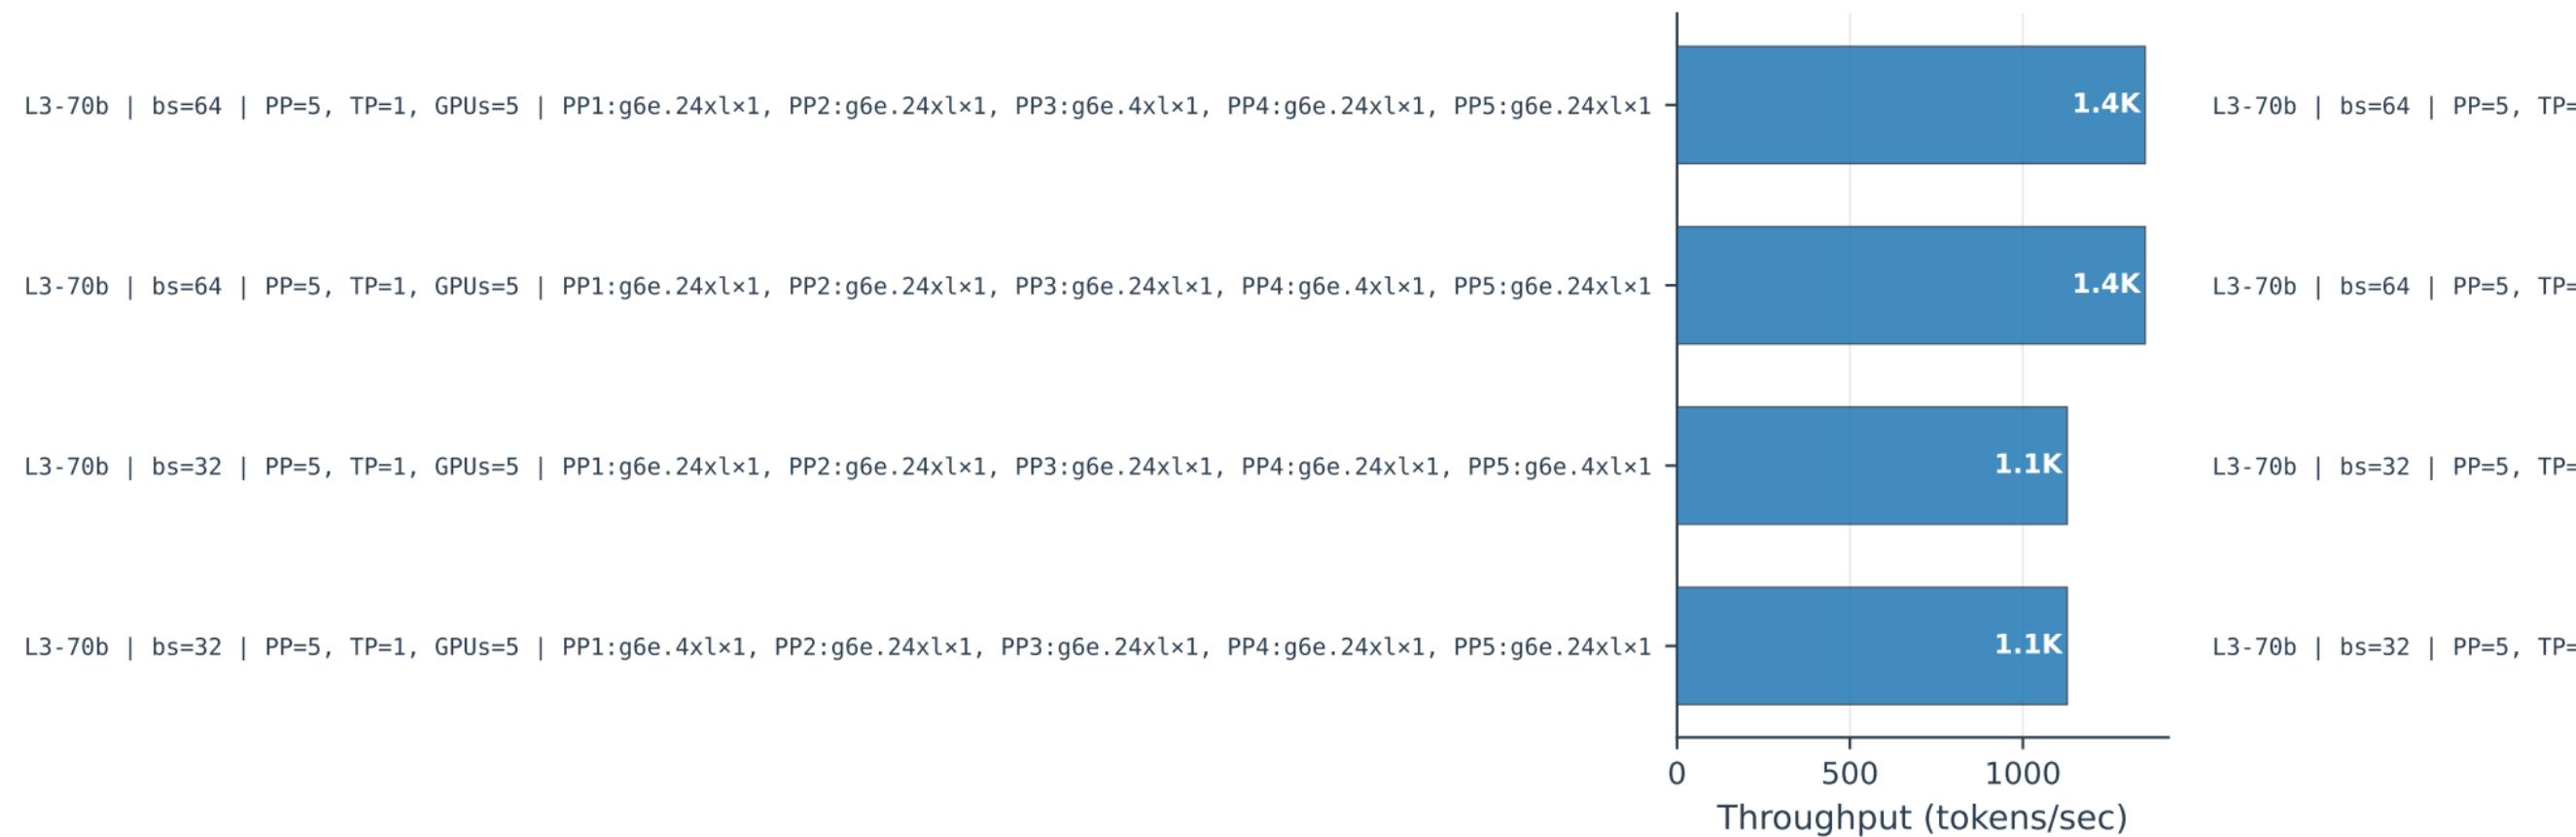
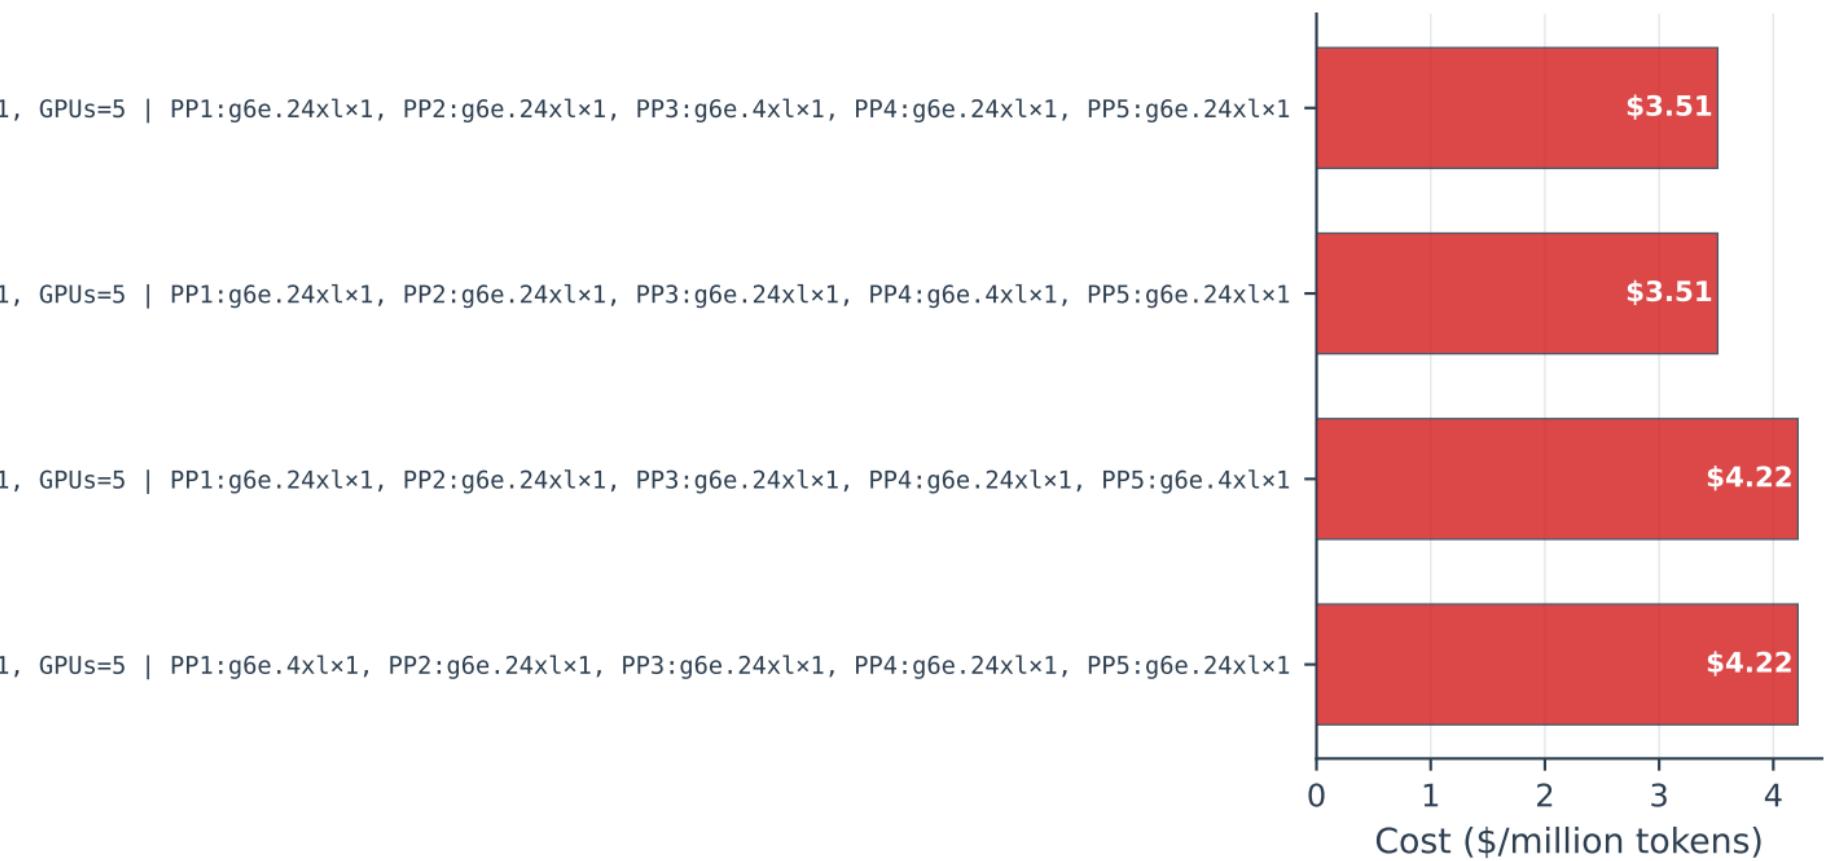


Prefill Workload — Input=1024, Output=1024 (Best Configurations Only)

Top 4 by Throughput

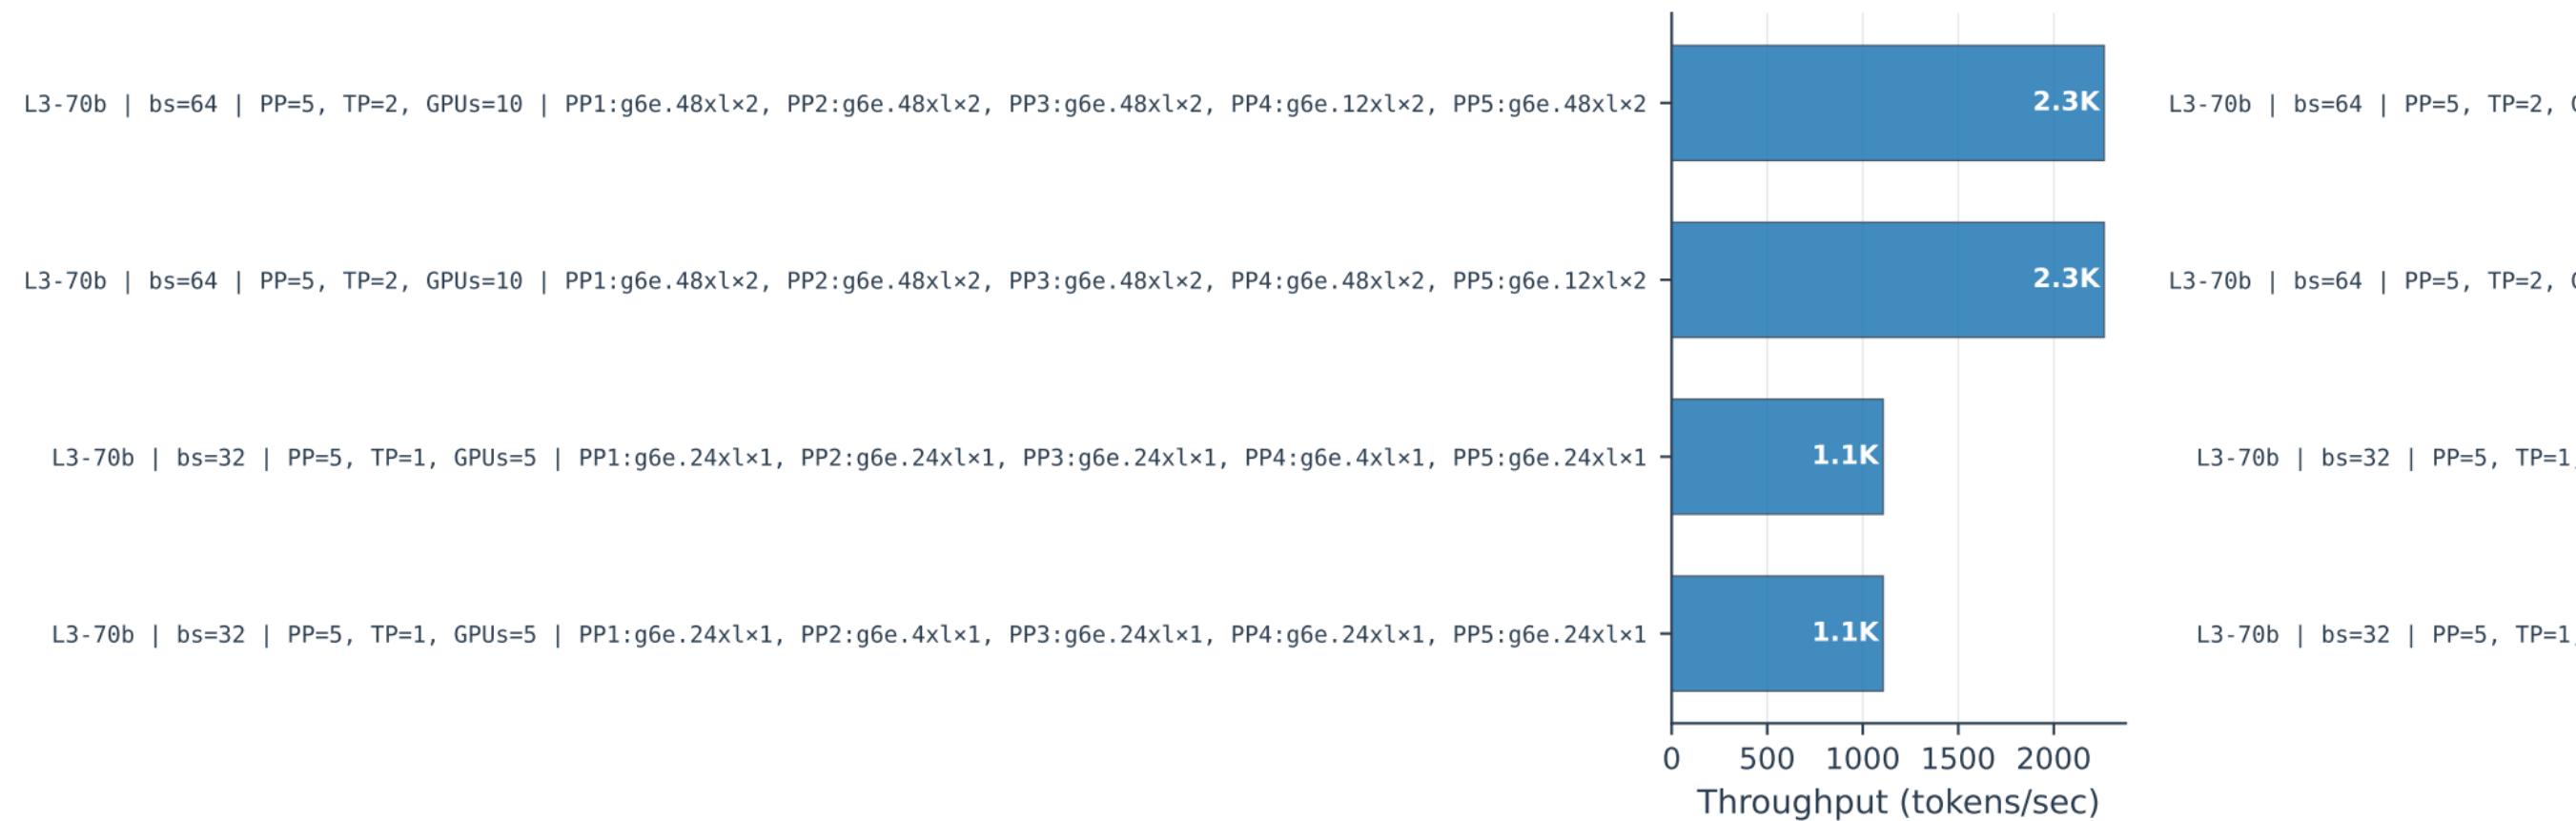


Top 4 by Cost Efficiency

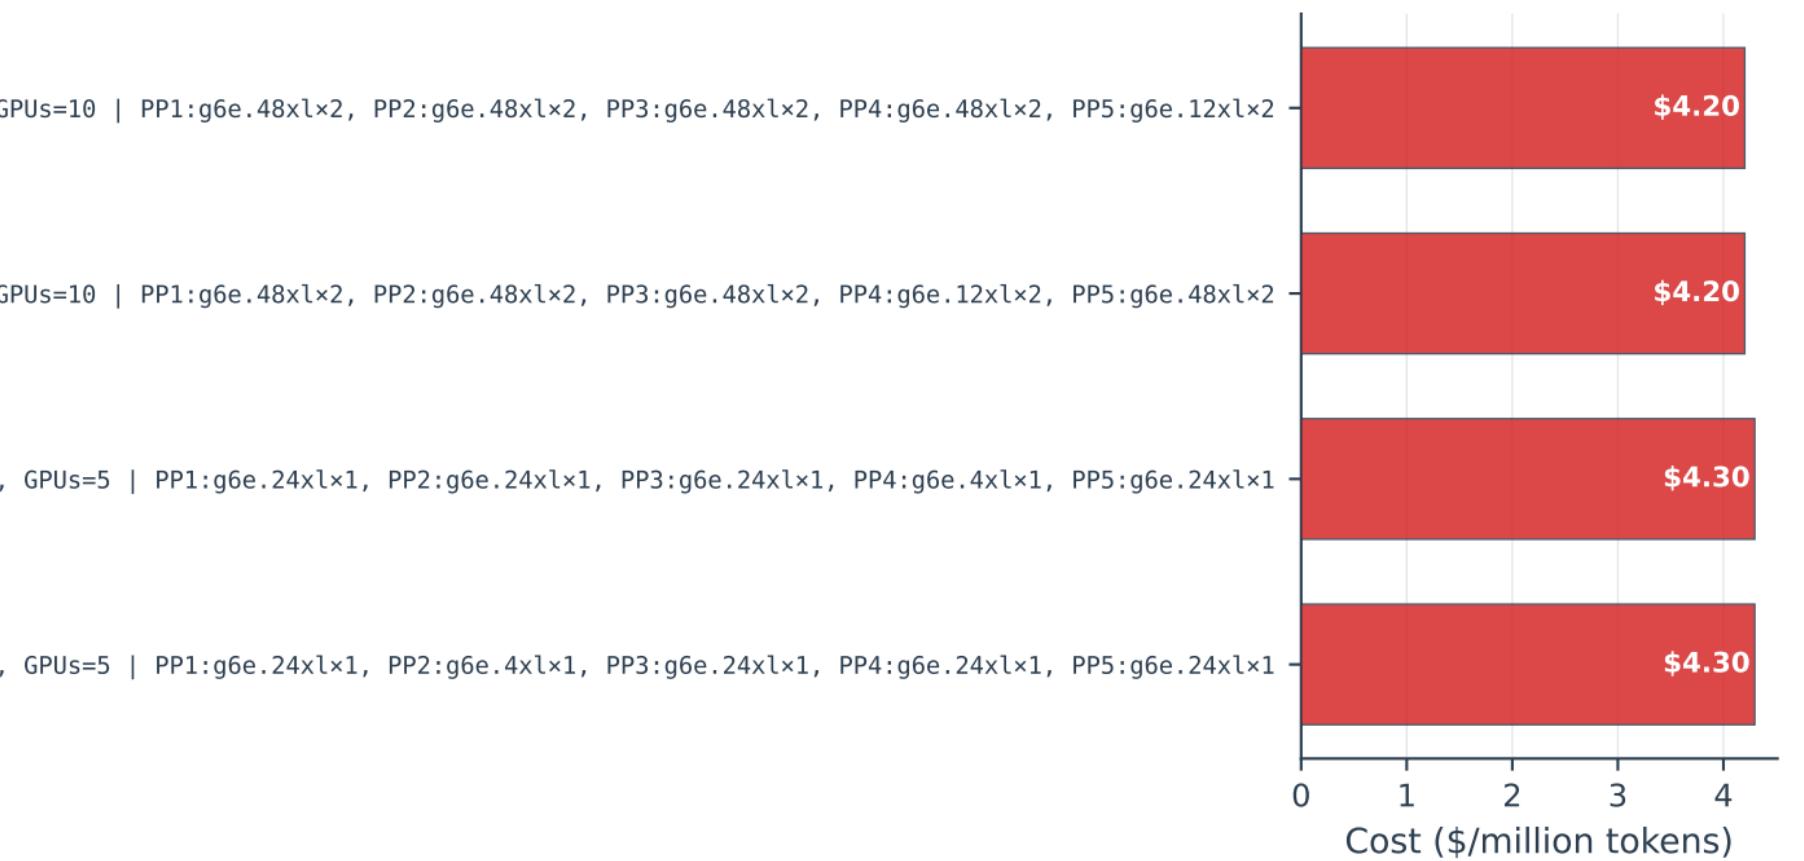


Prefill Workload — Input=2048, Output=2048 (Best Configurations Only)

Top 4 by Throughput

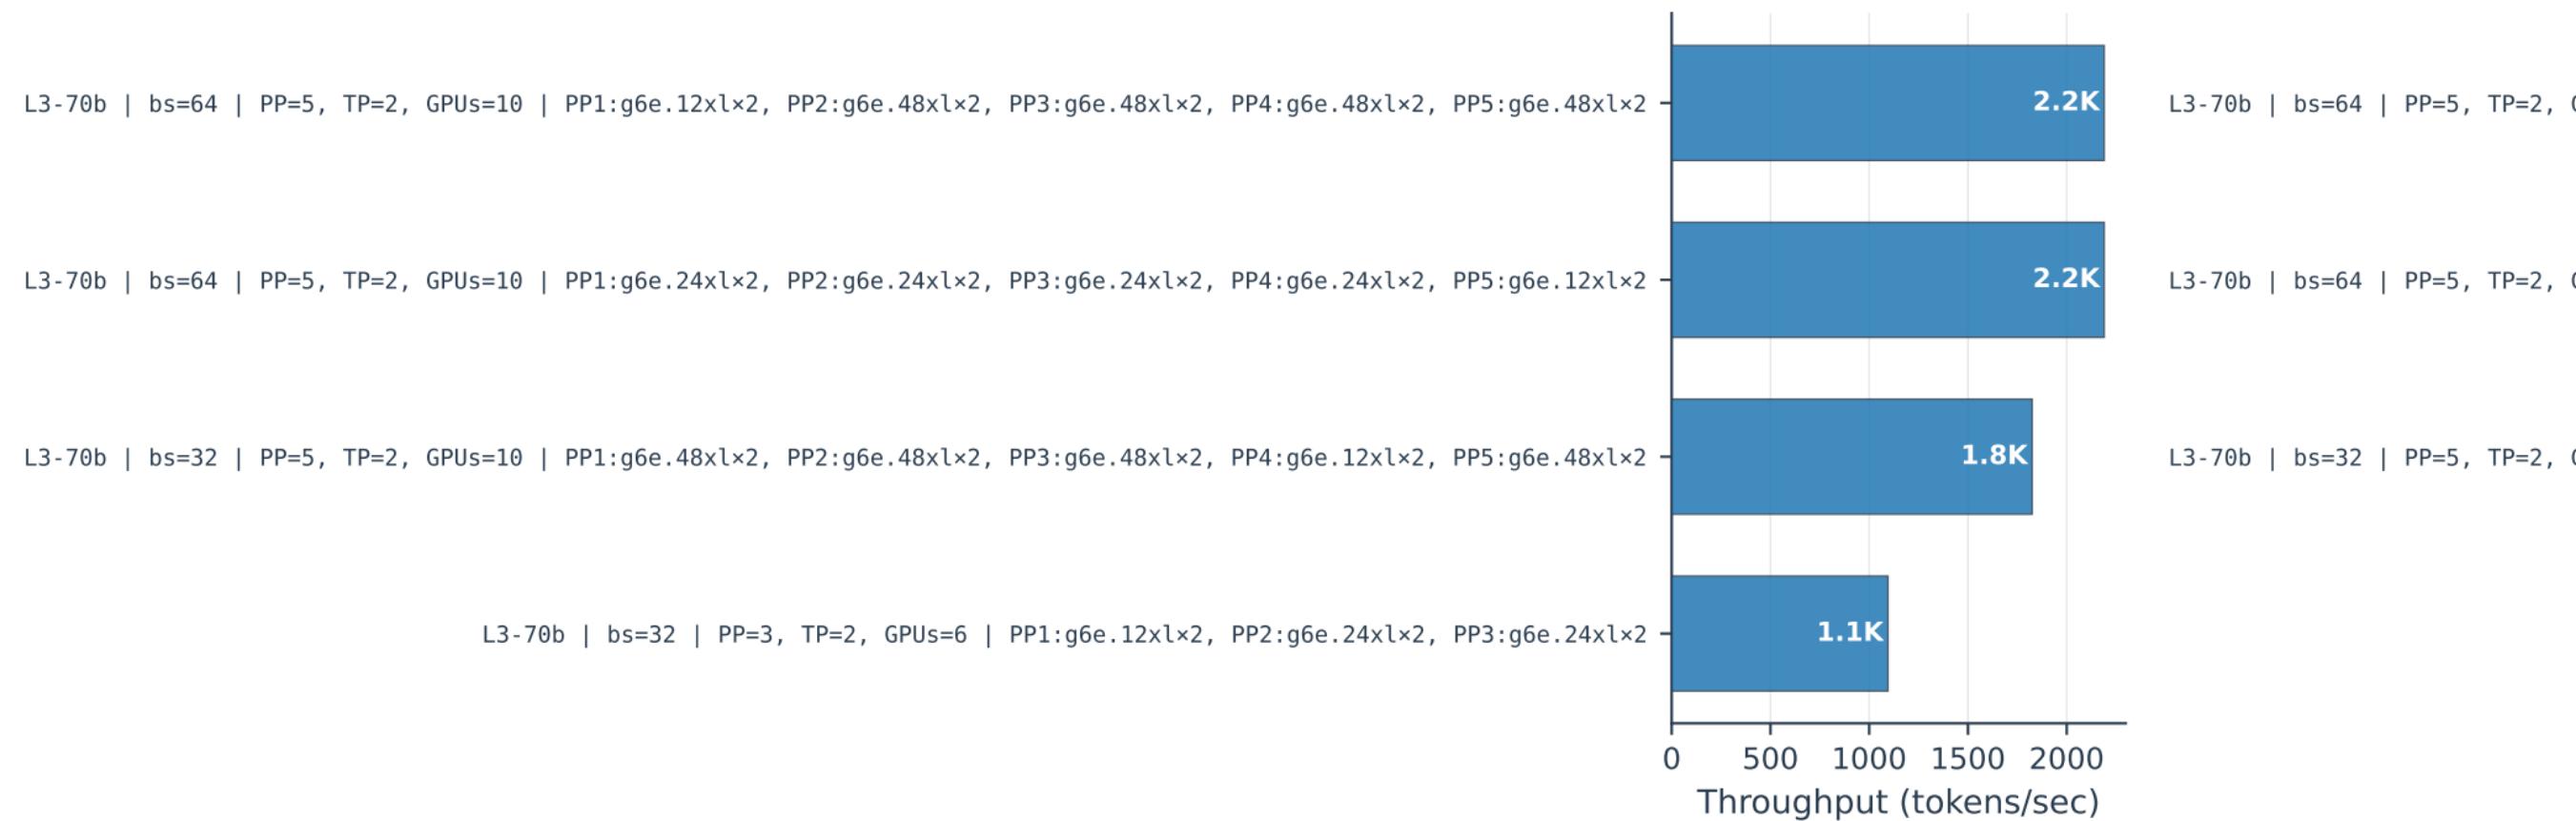


Top 4 by Cost Efficiency

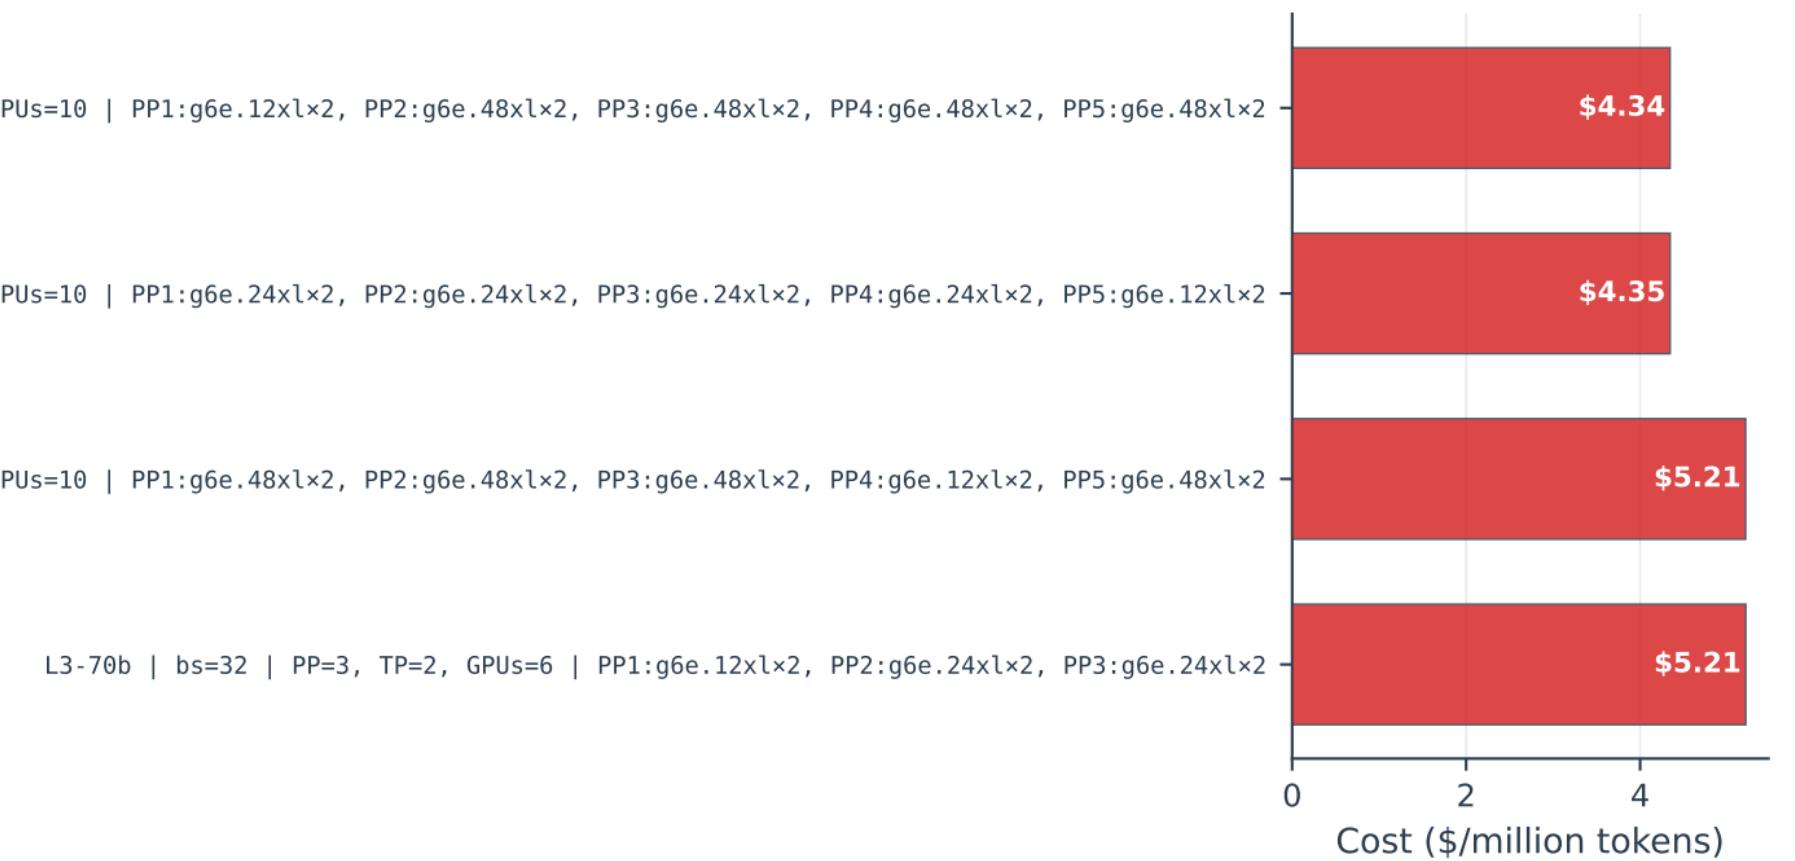


Prefill Workload — Input=4096, Output=4096 (Best Configurations Only)

Top 4 by Throughput

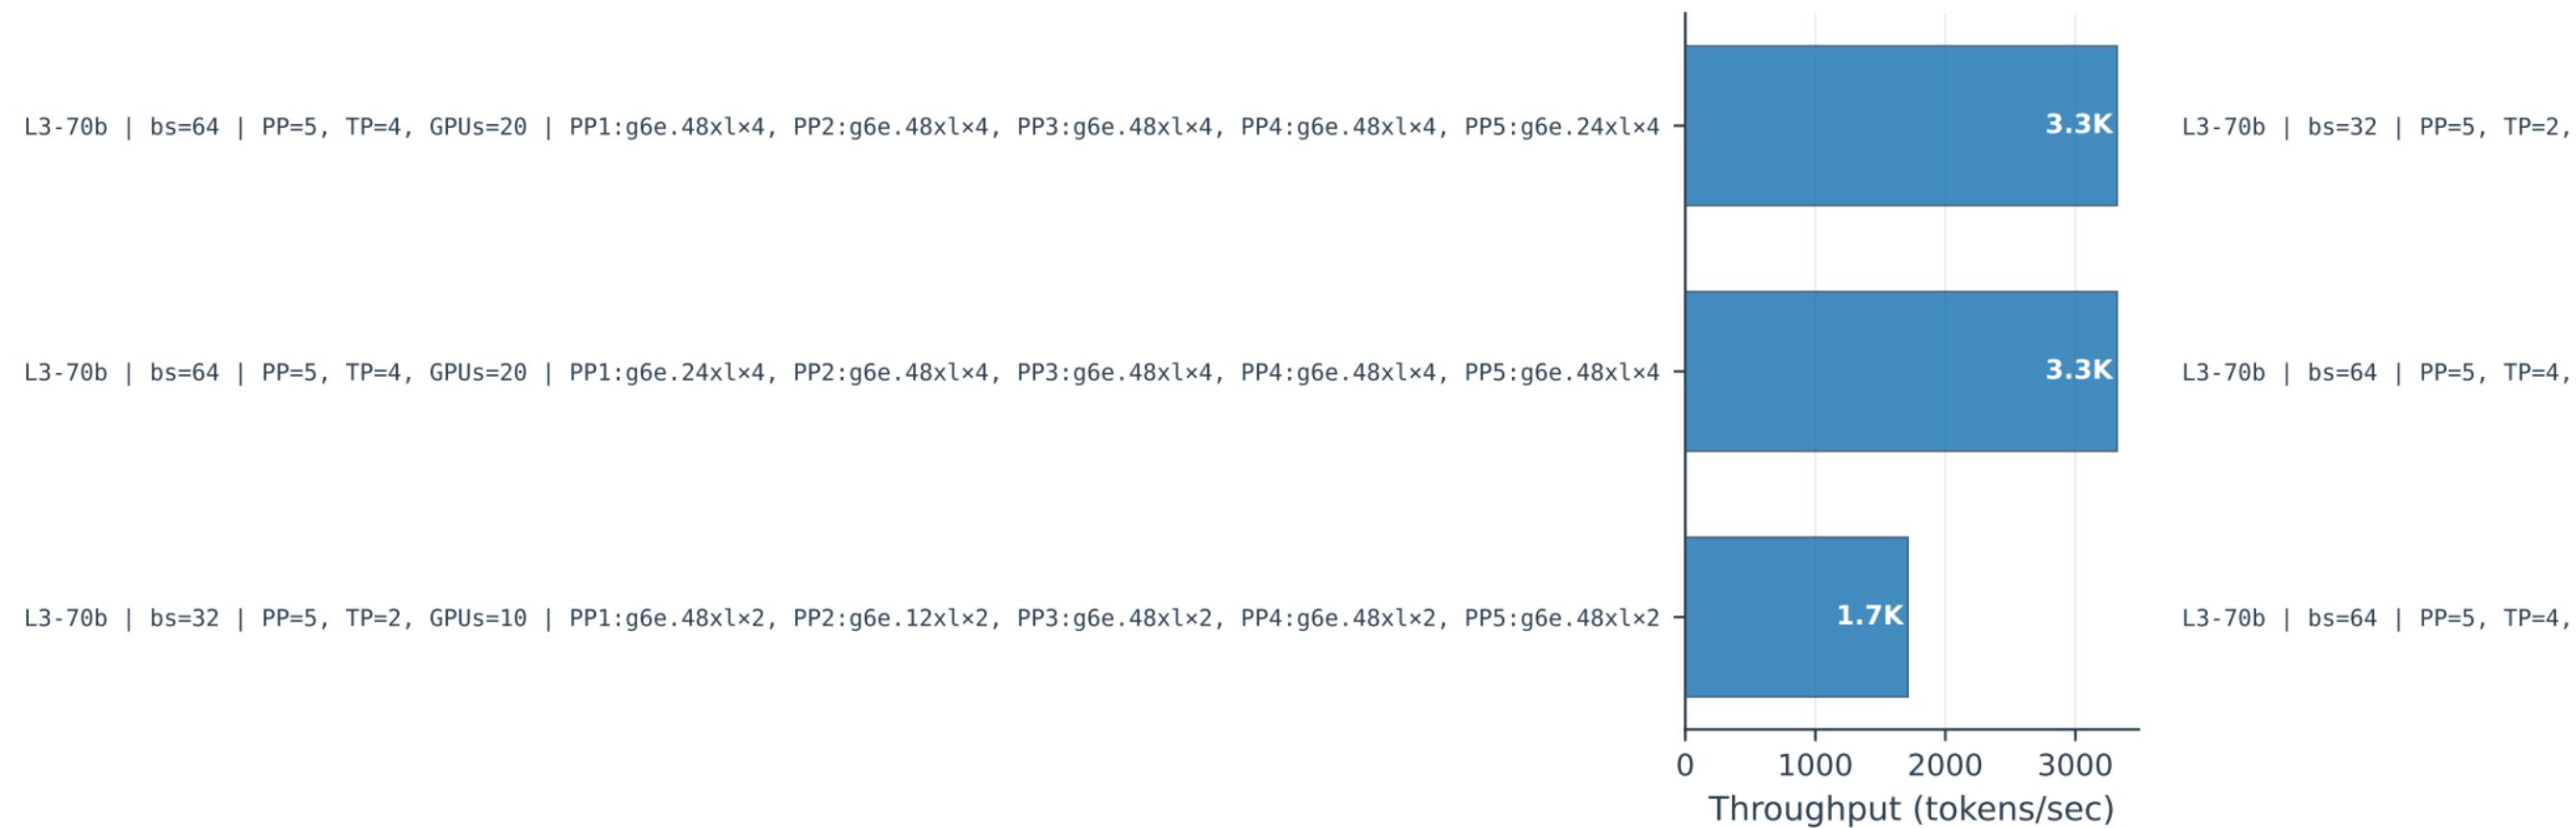


Top 4 by Cost Efficiency

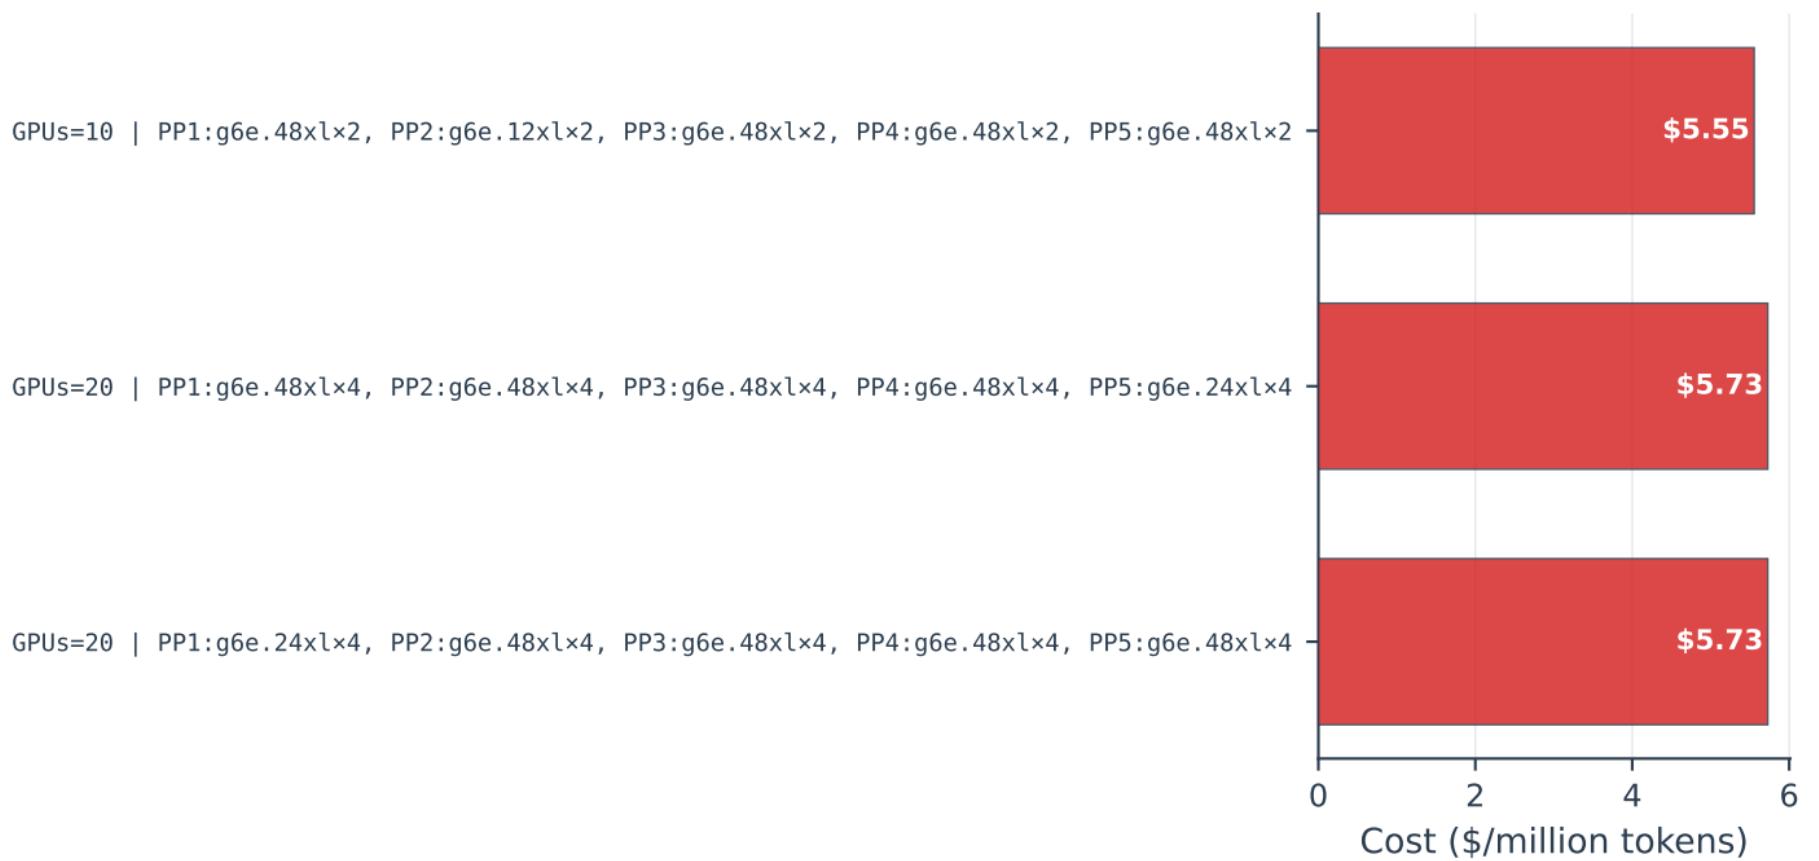


Prefill Workload — Input=8192, Output=8192 (Best Configurations Only)

Top 3 by Throughput

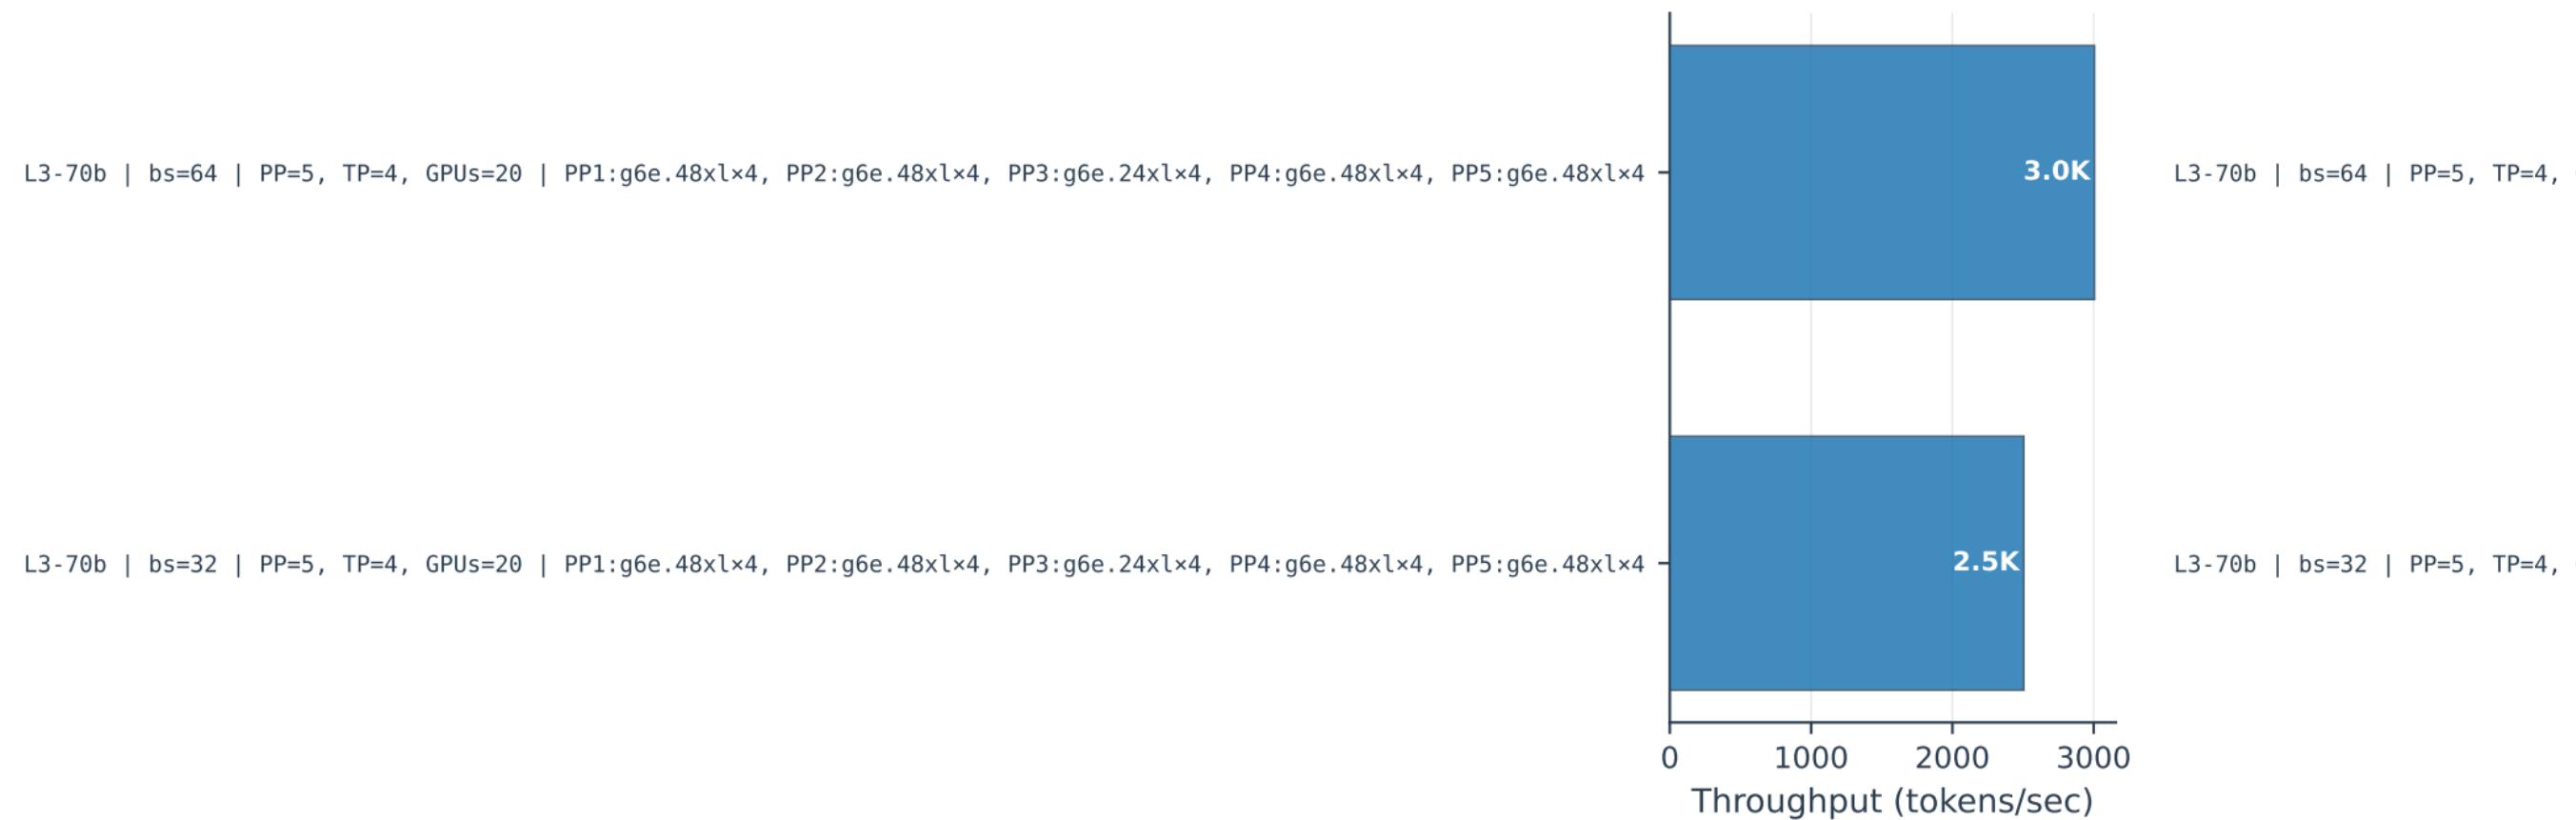


Top 3 by Cost Efficiency

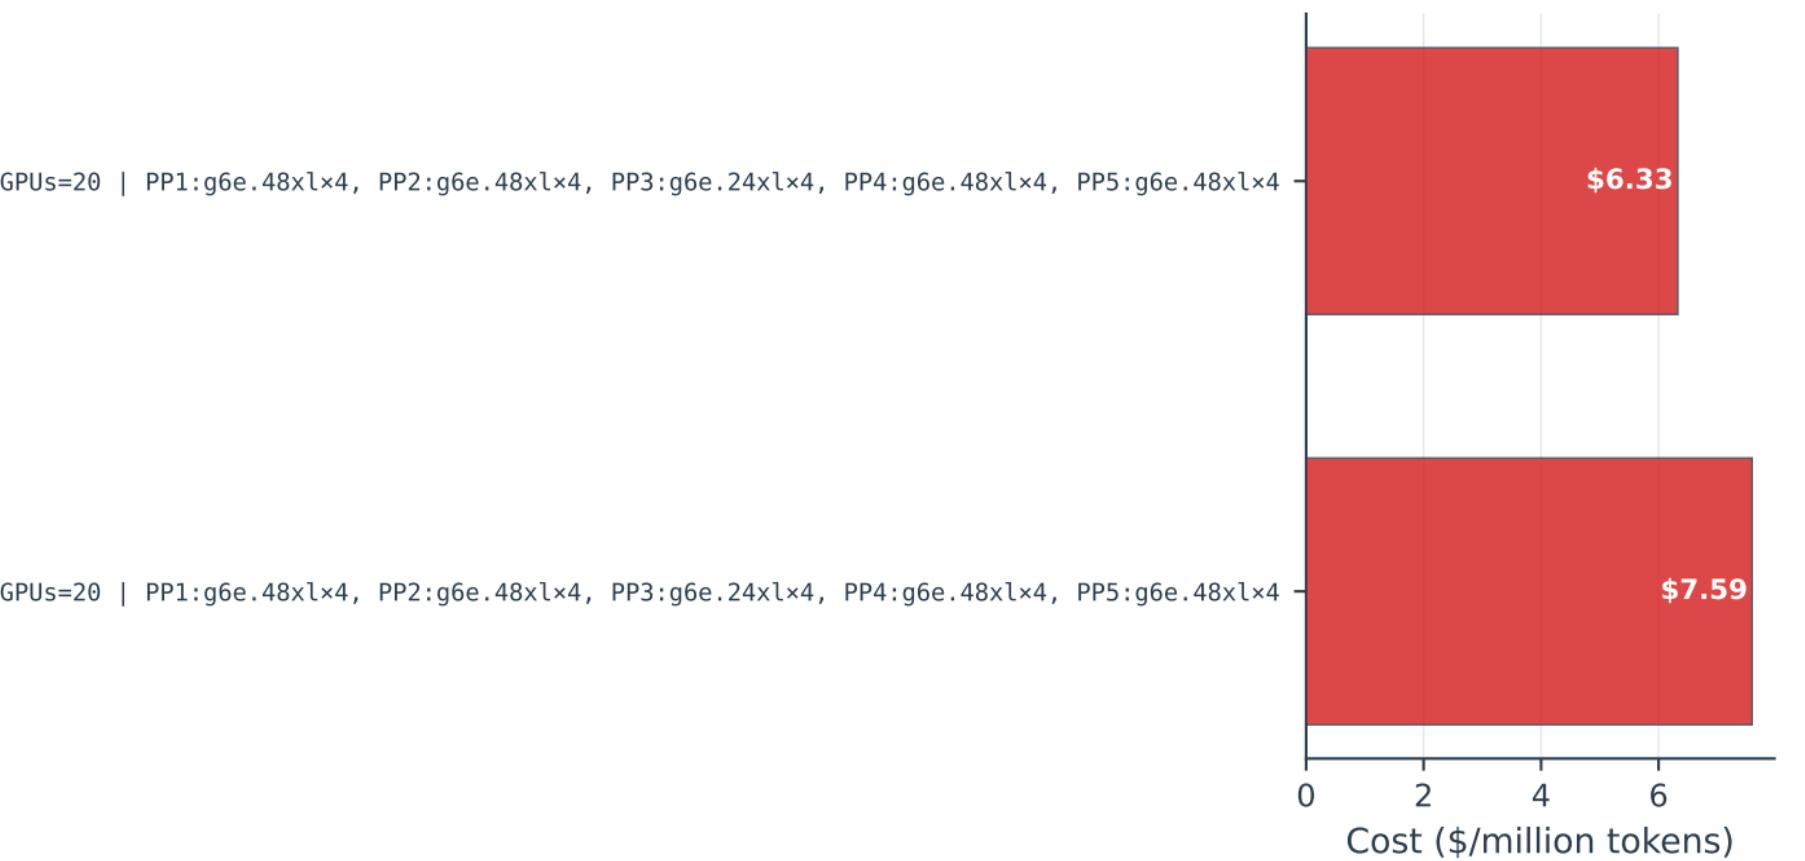


Prefill Workload — Input=16384, Output=16384 (Best Configurations Only)

Top 2 by Throughput

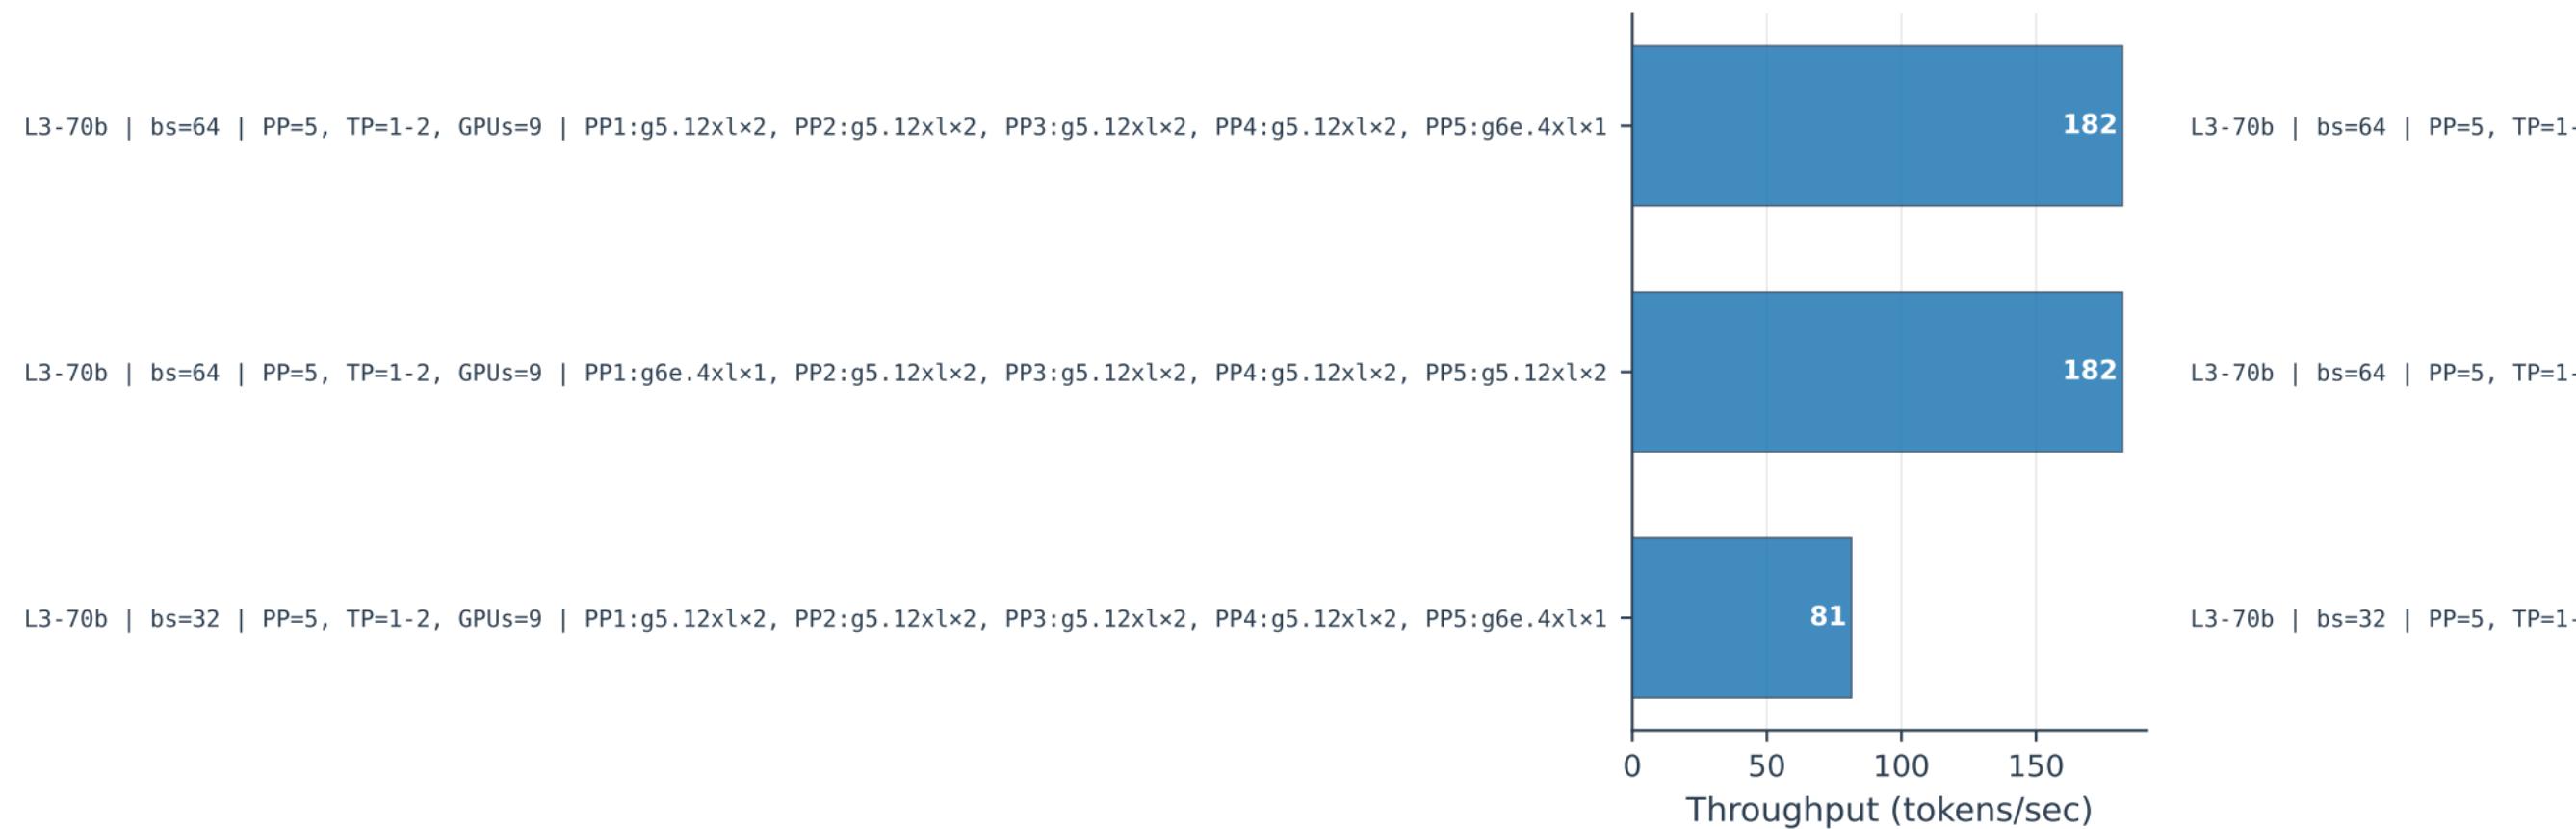


Top 2 by Cost Efficiency

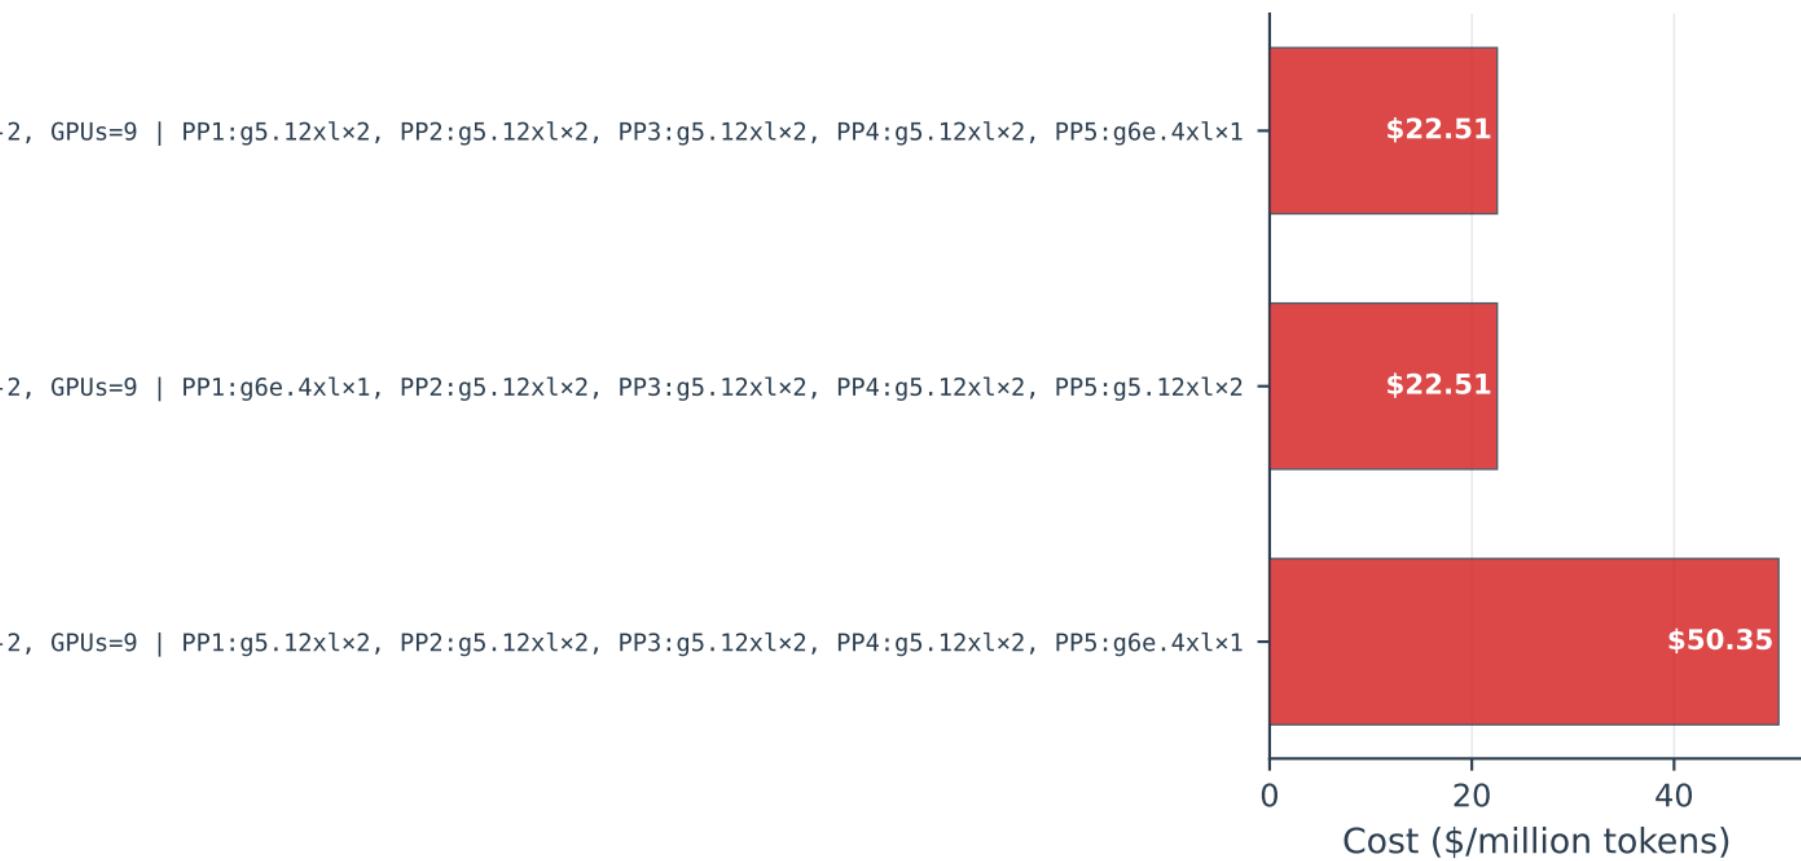


Decode Workload — Input=1024, Output=1024 (Best Configurations Only)

Top 3 by Throughput

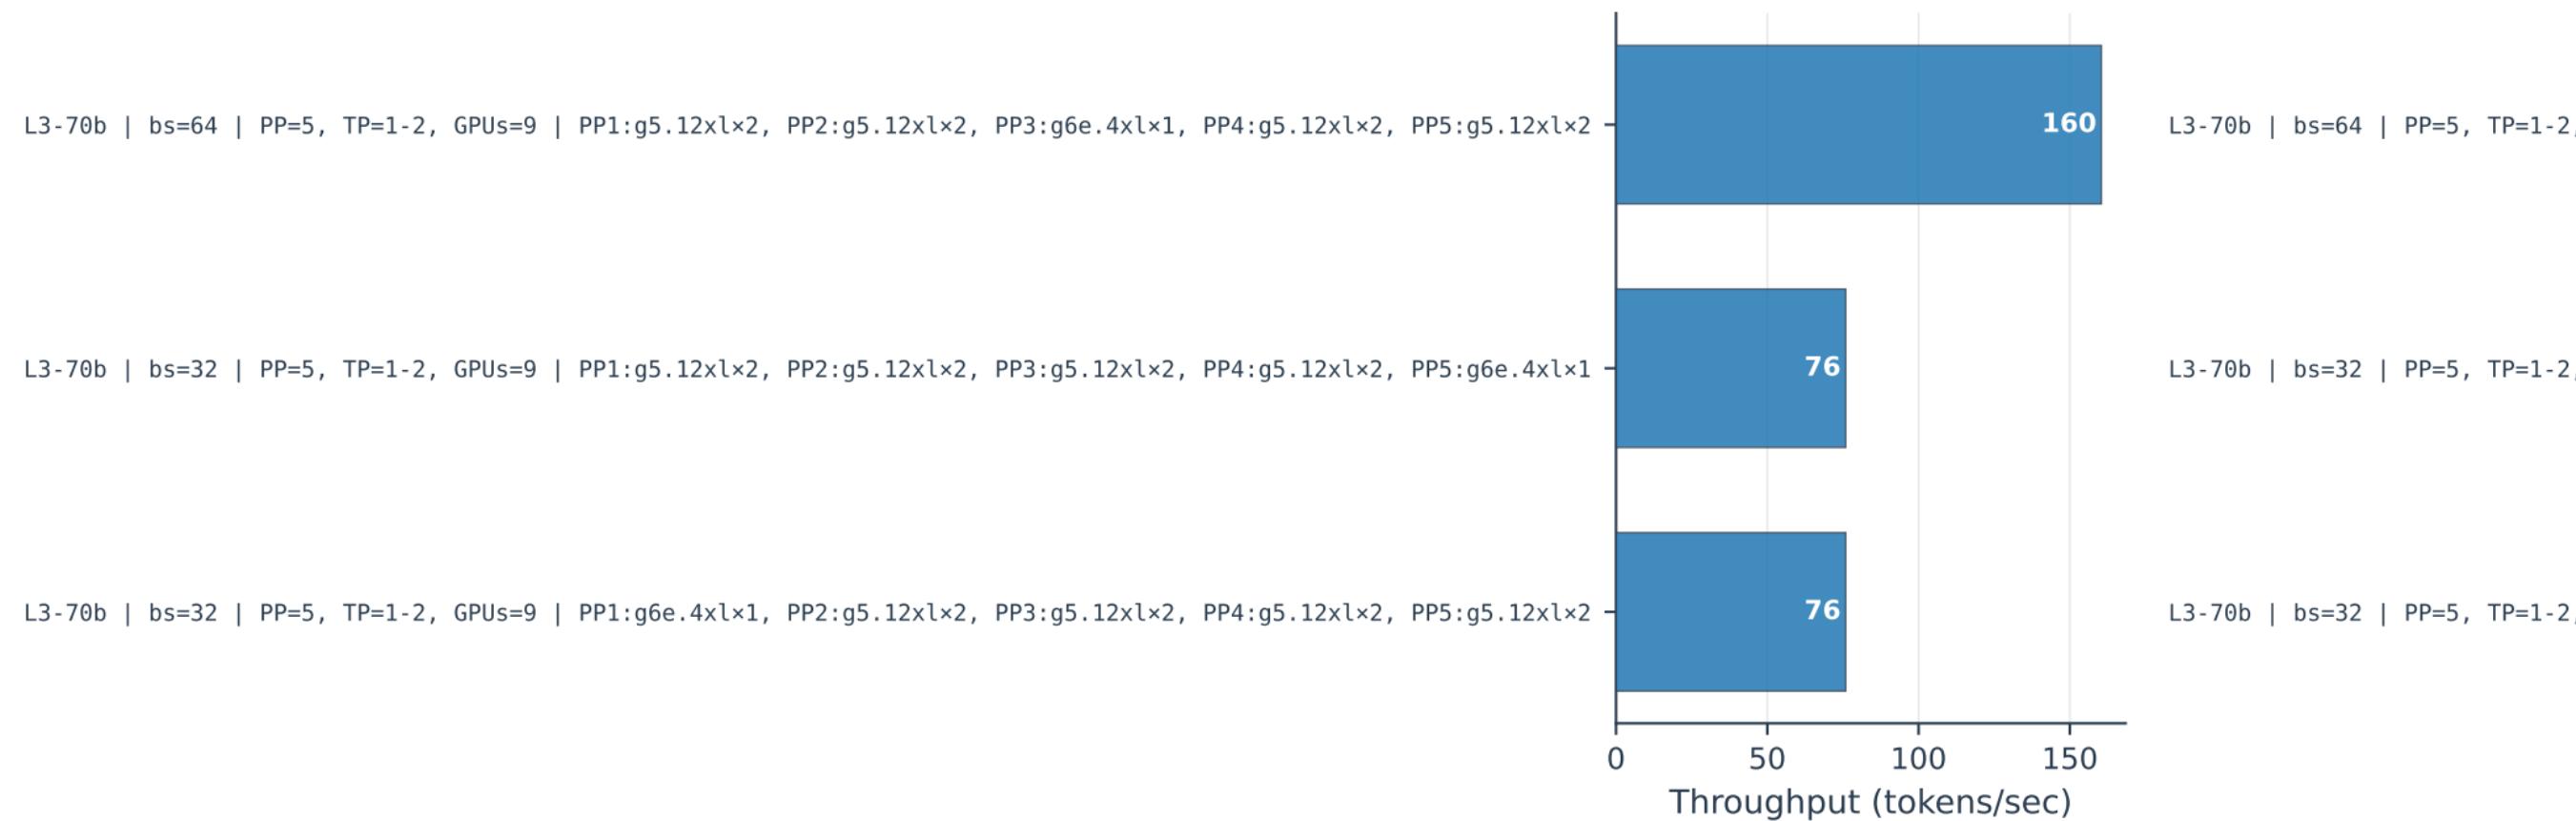


Top 3 by Cost Efficiency

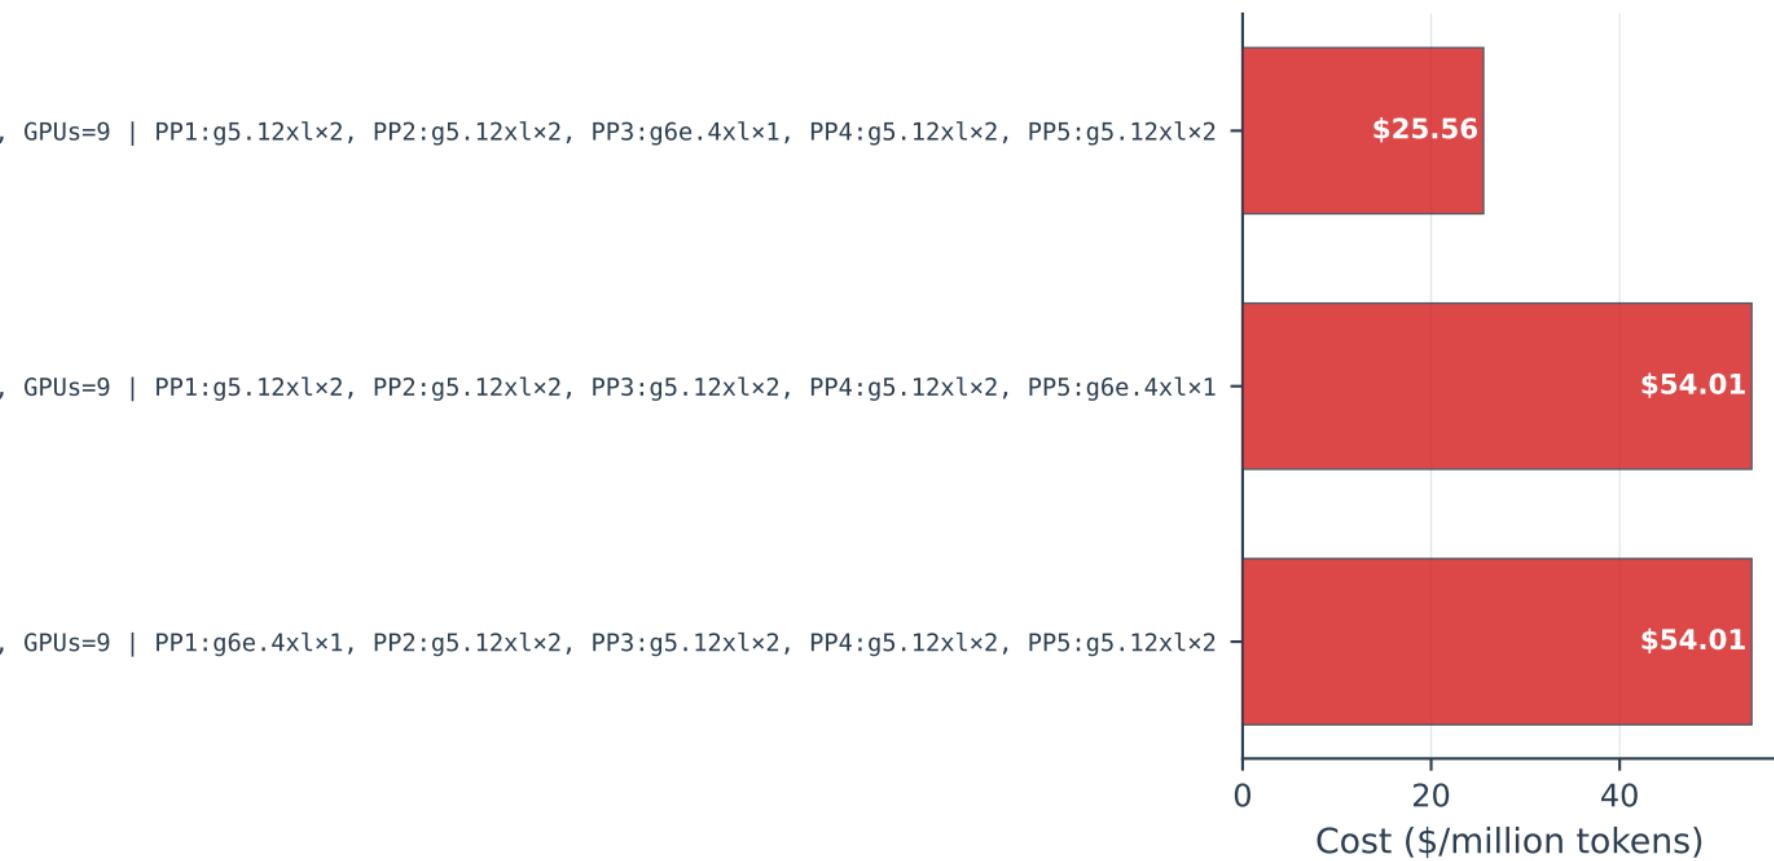


Decode Workload — Input=2048, Output=2048 (Best Configurations Only)

Top 3 by Throughput

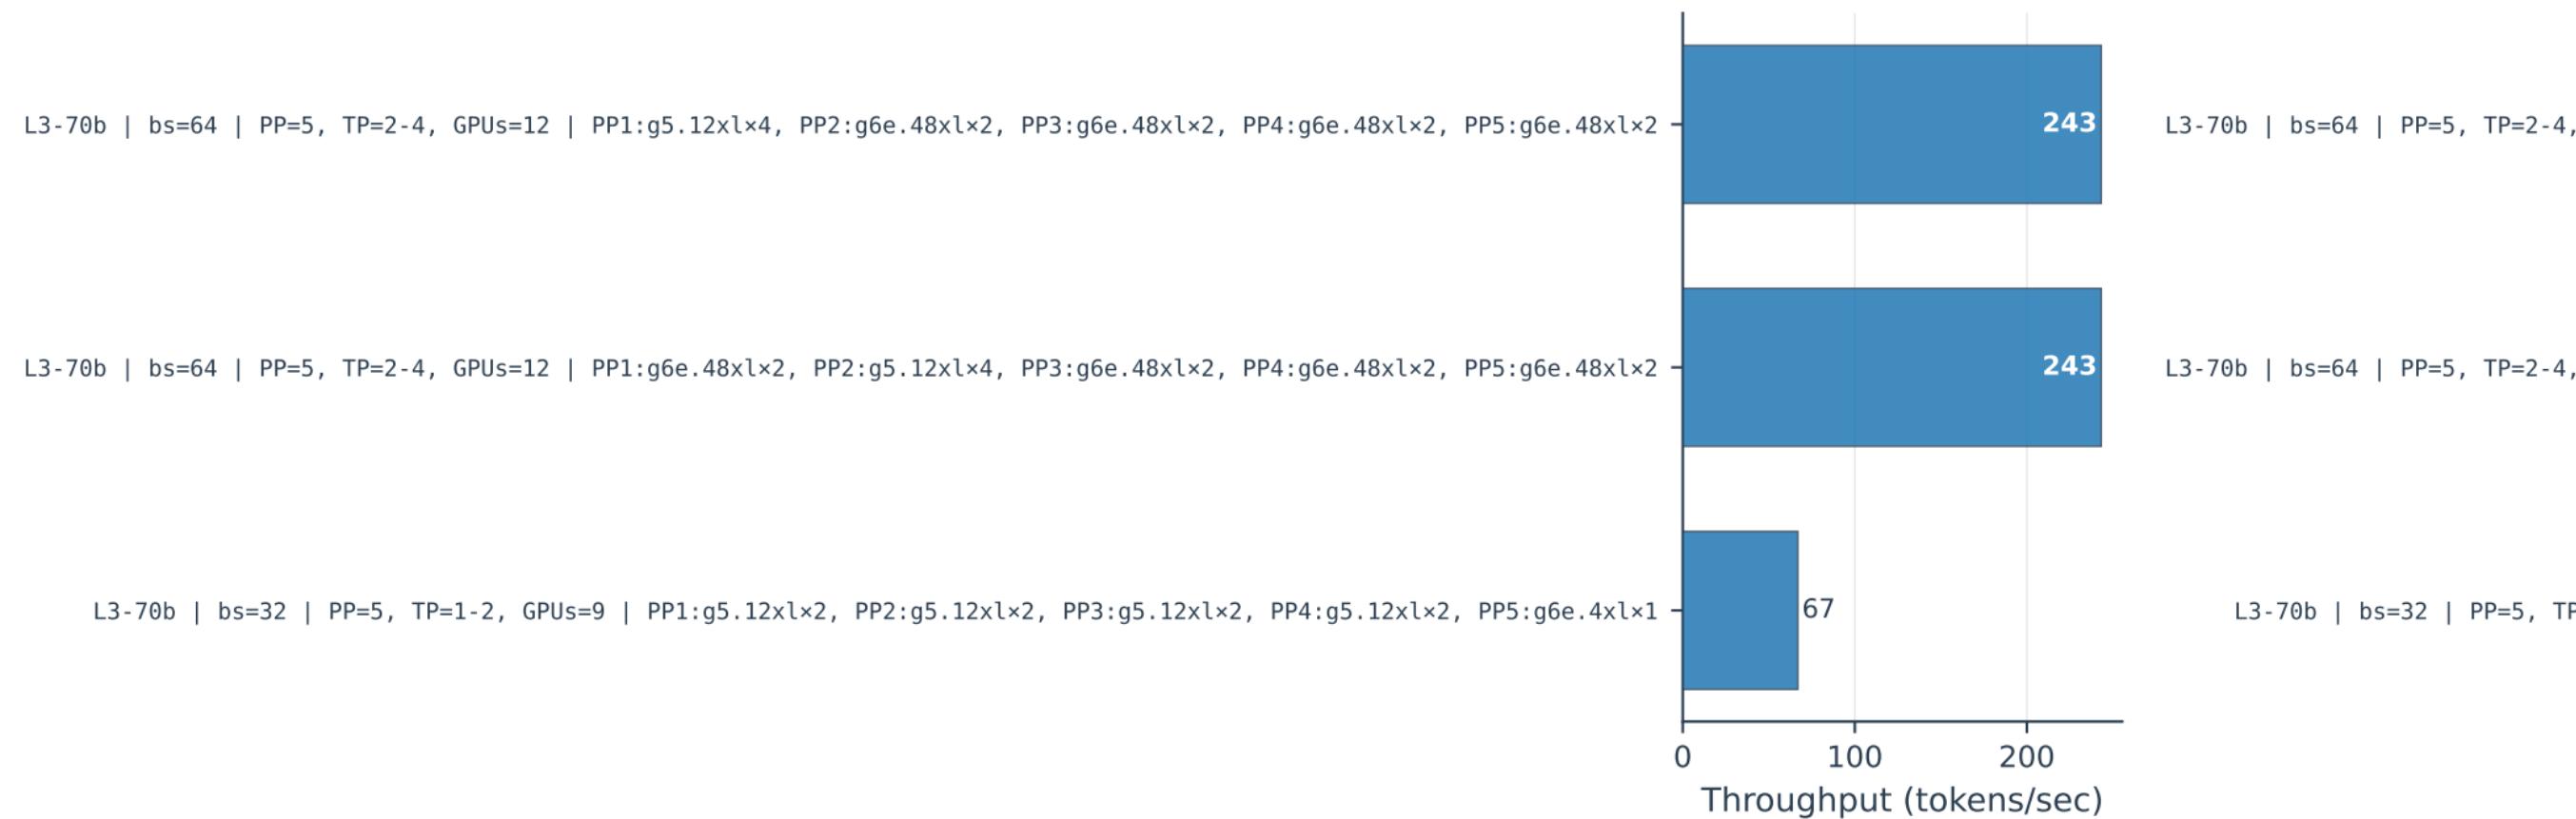


Top 3 by Cost Efficiency

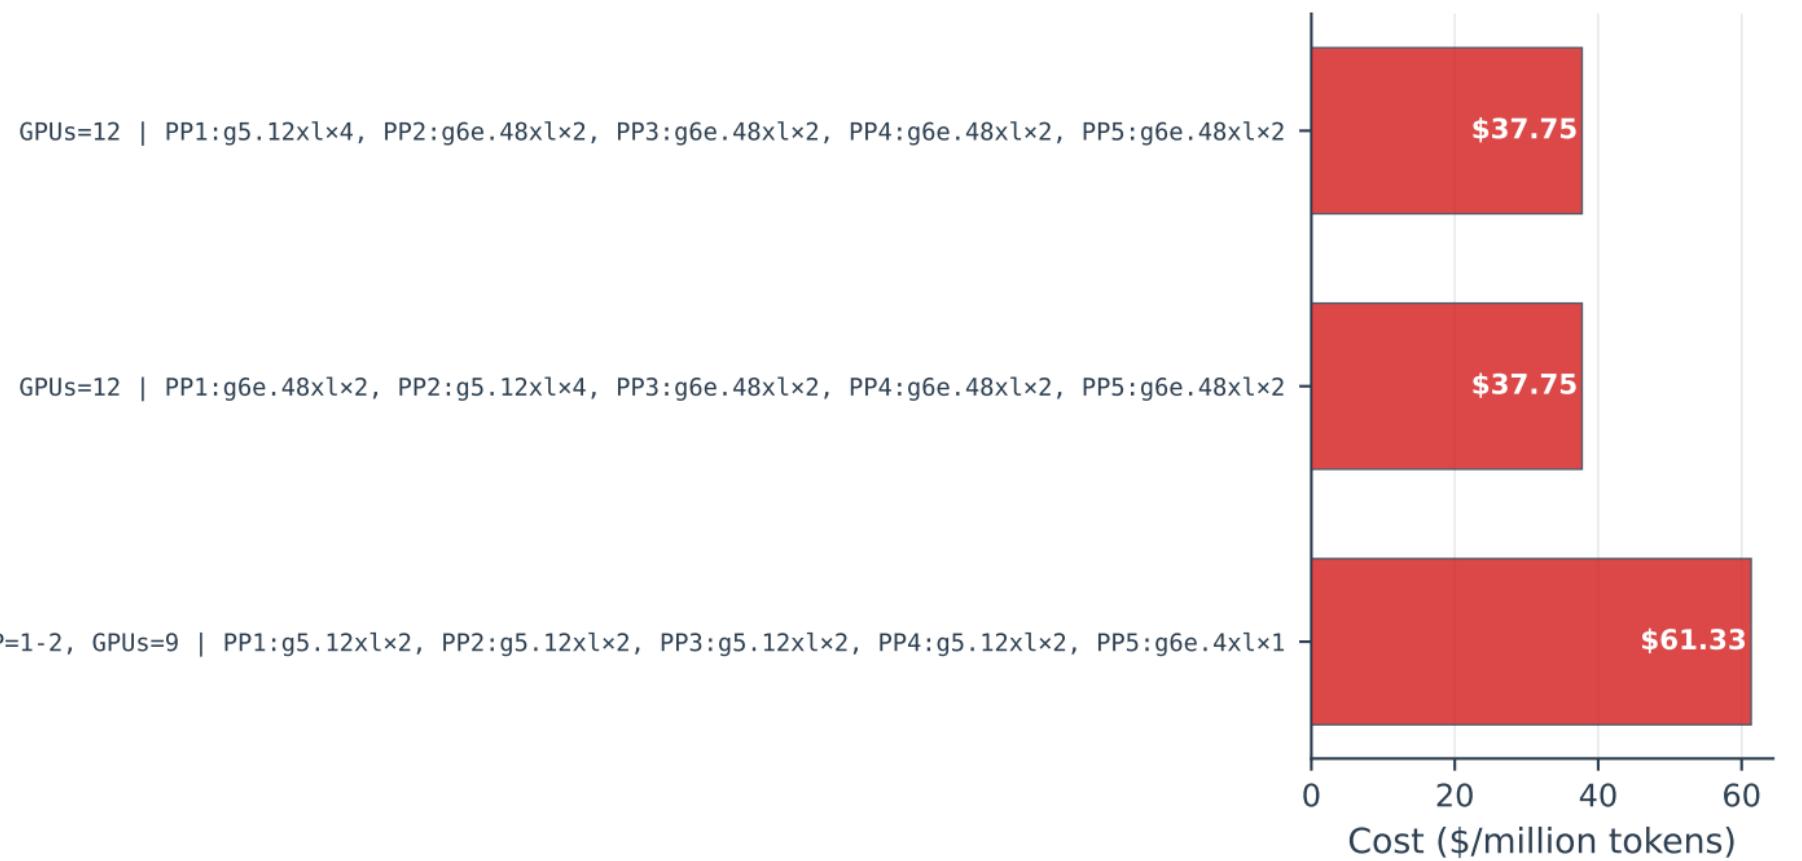


Decode Workload — Input=4096, Output=4096 (Best Configurations Only)

Top 3 by Throughput

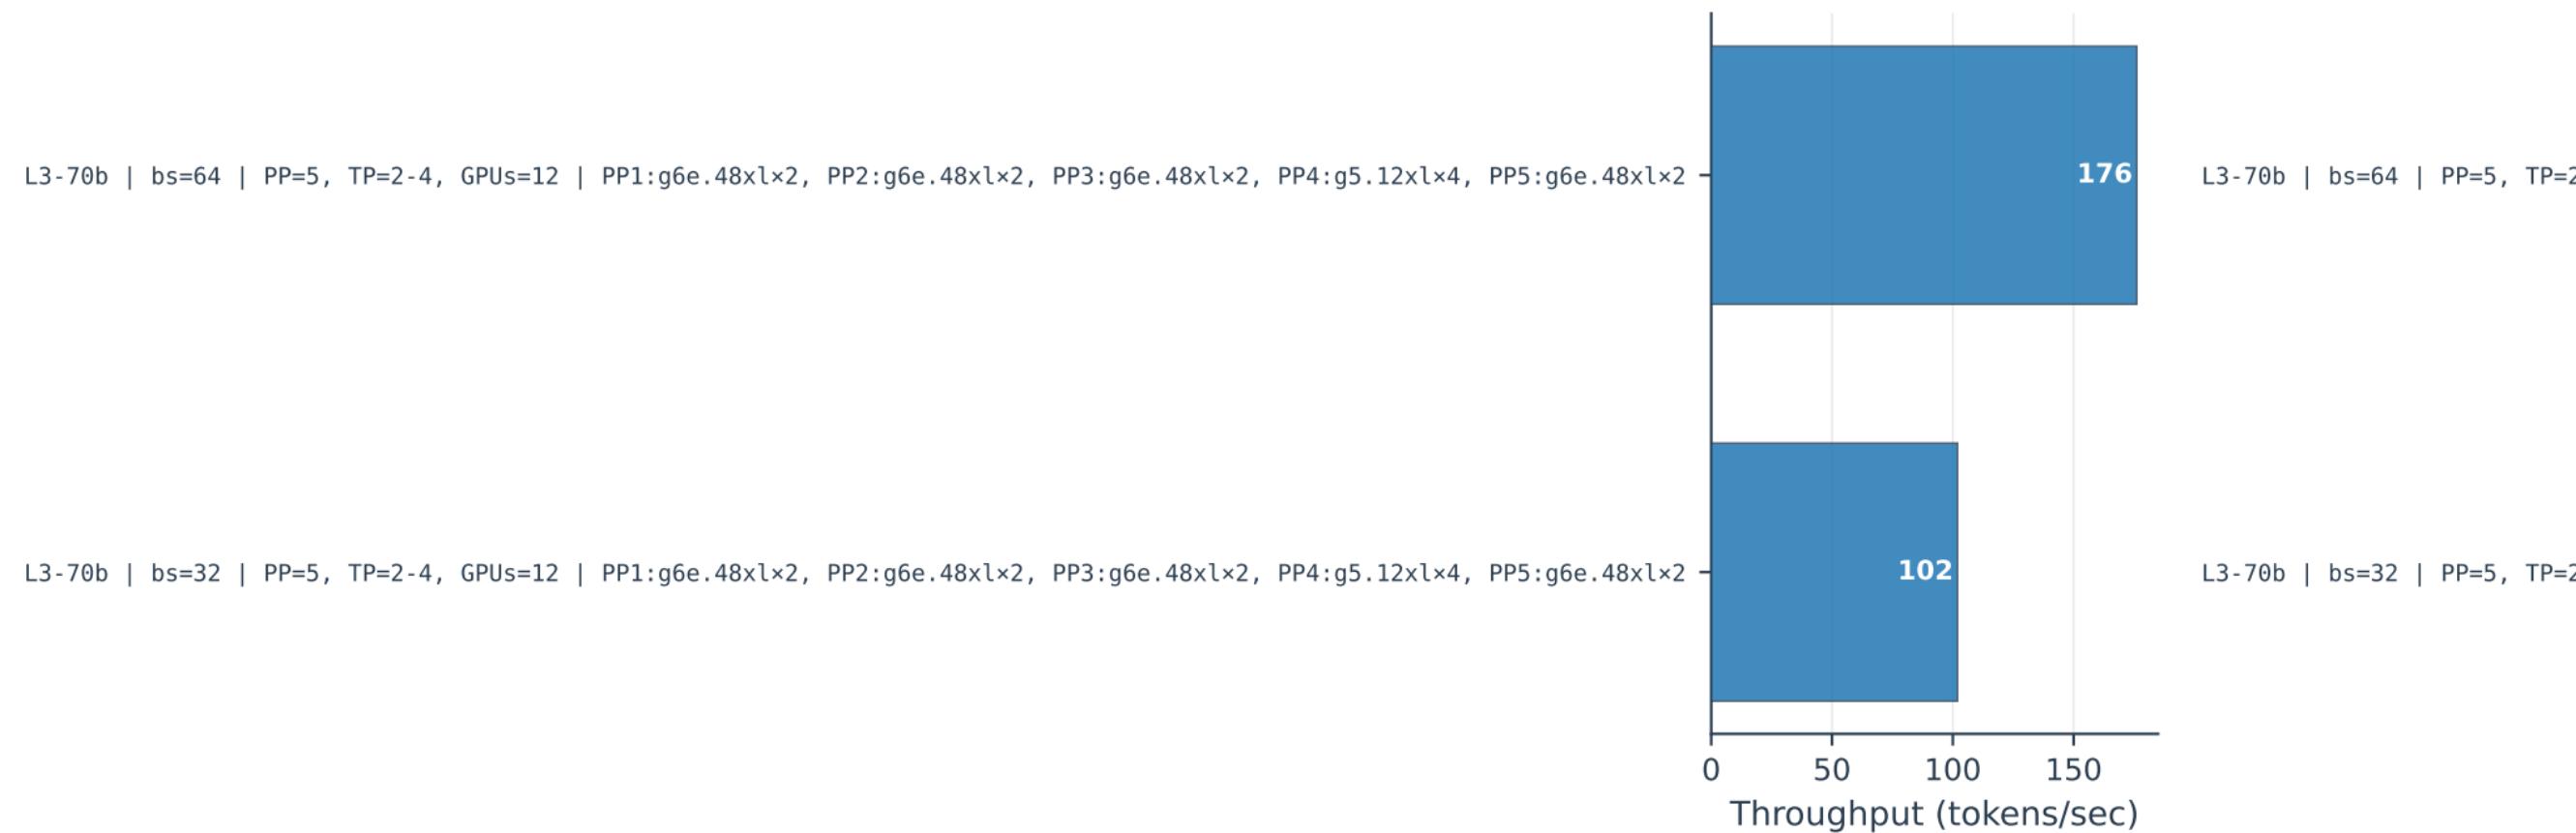


Top 3 by Cost Efficiency

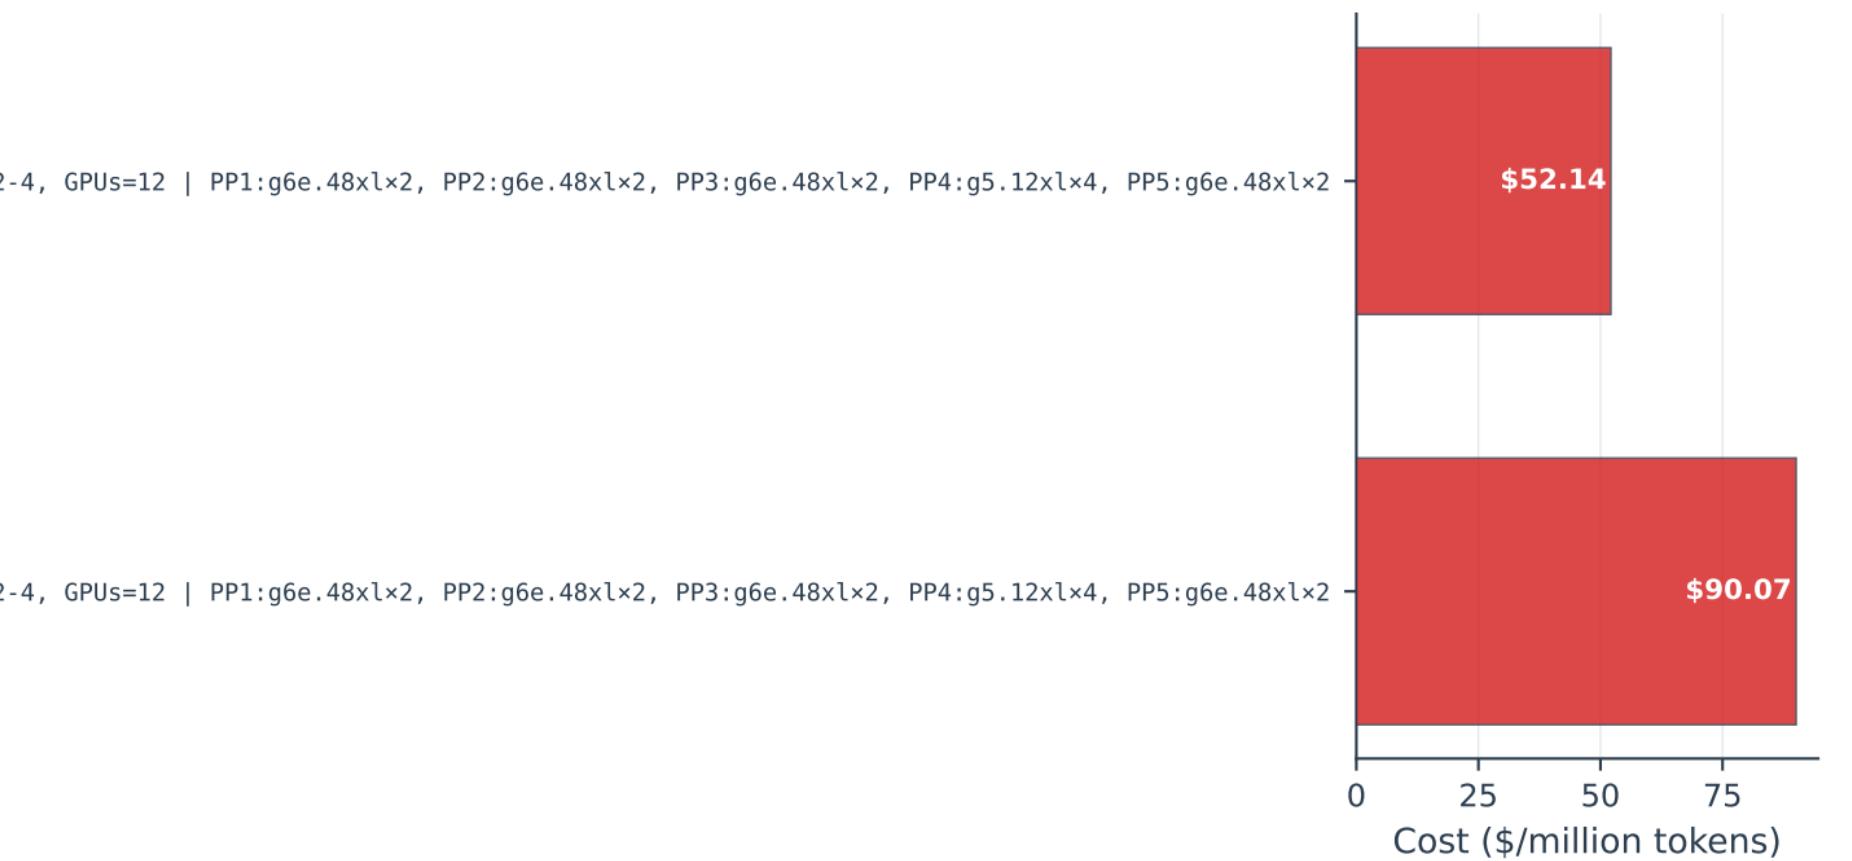


Decode Workload — Input=8192, Output=8192 (Best Configurations Only)

Top 2 by Throughput

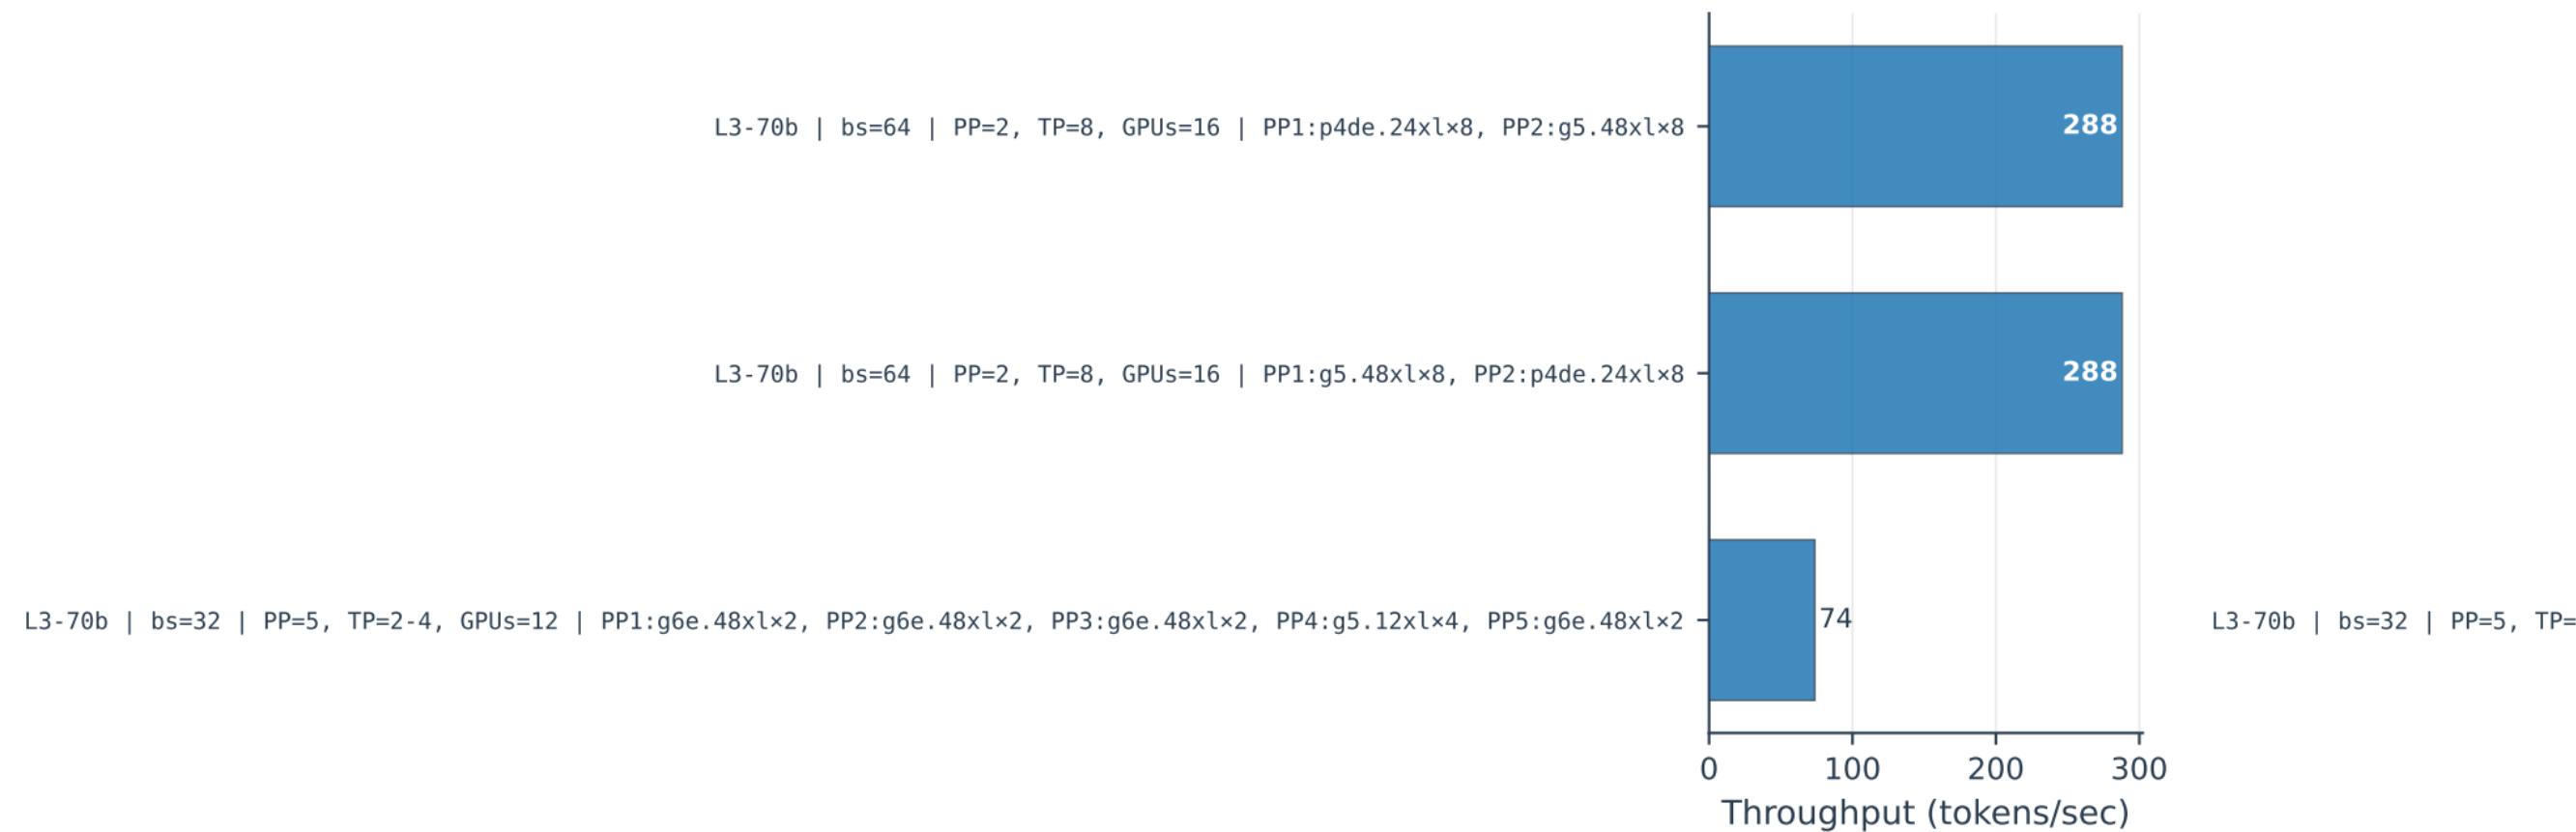


Top 2 by Cost Efficiency



Decode Workload — Input=16384, Output=16384 (Best Configurations Only)

Top 3 by Throughput



Top 3 by Cost Efficiency

